



CARB-v4 Master Benchmark Overview: Terminal-Bench 2.1 (89 Tasks • Gemini 3.6 Flash High)

(A) Macro Accuracy (N=89 Tasks)




(B) Wall-Clock Latency (Lower is Better)



(C) Unit Economics (Lower is Better)



(D) Hard Task Frontier (N=30)



Note: Evaluated across 267 isolated container sessions. Aggregate pass differences are not statistically significant (McNemar $p > 0.05$). Costs modeled at OpenRouter evaluation rates (0.5598in/3.745 out per M tokens).